◈ PROJECT OVERVIEW ◈

JiminAI Lie Detector *PATENT PENDING*

"Alignment isn't a policy. It's a measurement."

WHAT YOU'RE SEEING

This live demonstration detects whether an AI model is being deceptive — in real time. We never look at the tokens. We never read the output.

No peeking. No tricks. Just math.

Think of it like reading body language — but for AI. The model can say whatever it wants — we're reading its brain waves, not its words.

HOW IT WORKS

Every time a prompt is processed, we:

1. Run the prompt through Qwen2.5-14B-Instruct
2. Capture the model's internal state during inference
3. Run the math
4. Deliver a verdict: HONEST or DECEPTIVE — before you even finish reading the response

UNDER THE HOOD

Model: Qwen2.5-14B-Instruct (28GB, float16 precision)

Server: 3x NVIDIA RTX 3090 GPUs (24GB each), 125GB RAM, running on bare metal in Austin, TX

Inference: Direct HuggingFace transformers — no API calls, no cloud. The model runs entirely on our hardware. Every KV cache is captured in-process during generation.

Detection: Mathematical analysis of internal model state. Research-validated AUROC of **0.983** across thousands of test inferences.